

SQL: A COMMERCIAL DATABASE LANGUAGE

Data Definition,
Basic Integrity Constraints,
Schema Changes

Outline of Chapters 8, 9

1. Introduction
2. Data Definition, Basic Constraints, and Schema Changes
3. Basic Queries
4. More complex Queries
5. Aggregate Functions and Grouping
6. Summary of SQL queries
7. Data Change statements
8. Views
9. Complex Constraints
10. Database Programming

1. Introduction (1)

- The name **SQL** is derived from **S**tructured **Q**uery **L**anguage. Originally SQL was called **SEQUEL** (for Structured English QUery Language).
- SQL is a *comprehensive database language*. It has statements for *data definition, query* and *update* (in this sense it is both a **DDL** and a **DML**). In addition, it has facilities for *defining views*, for specifying *security and authorization*, for defining *integrity constraints*, and for specifying *transaction controls*. It also has rules for *embedding SQL statements* into a general-purpose programming language.
- SQL became a *standard for relational databases* in 1986 (called **SQL-86** or **SQL1**). A revised and more extended version became a standard in 1992 (called **SQL-92** or **SQL2**). A new version of the standard (called **SQL:1999**) extends SQL with object-oriented and other recent database concepts.

1. Introduction (2)

- The SQL language provides a *high-level declarative language* interface: the user specifies what the result is to be and leaves the decisions on how to execute the query to the DBMS.
- SQL includes features from *Relational Algebra* but it is based to a greater extend on *Tuple Relational Calculus*. However, Relational Algebra is very important for *query processing and optimization* of SQL queries in relational DBMS.
- In class we will overview the most important features of SQL. Details can be found in SQL documents.

1. Introduction (3)

- The terms **table**, **row** and **column** are used in SQL documents for **relation**, **tuple** and **attribute** respectively.
- *SQL is case insensitive*. It treats upper- and lower-case letters as the same letter. Only inside quotes does SQL make a distinction between upper- and lower-case letters.

2. Data Definition, Constraints and Schema Changes in SQL

- In this section we discuss DDL commands in SQL.

2.1.Schema and Catalog concepts (1)

- The concept of a **schema** is used in SQL to group together constructs that belong to the same database application
- A schema includes **schema elements** such as *tables, constraints, views, users* and other constructs.
- An SQL schema is identified by a **schema name** and includes an **authorization identifier** to indicate the user who owns the schema. In addition, it can include **descriptors** for every element in the schema.
- The following statement creates a **schema/database** called COMPANY

CREATE DATABASE COMPANY;

2.1.Schema and Catalog concepts (2)

- The concept of a **catalog** is used in SQL to denote a named collection of schemas.
- A catalog always contains a special schema called **INFORMATION_SCHEMA**. This schema provides information on all the schemas in the catalog and all the element descriptors in these schemas.
- Some elements (e.g. domain definitions) *can be shared by schemas in the same catalog*.

Integrity constraints can be defined between relations *only if they exist in schemas of the same catalog*.

2.2. Data Types (1)

- The data types available for attribute values include numeric, character-string, bit-string, date and time.
- **Numeric:**
 - INTEGER** or **INT**: signed integer
 - Also **TINYINT**, **SMALLINT**, **MEDIUMINT**, **BIGINT**
 - FLOAT**: floating point number
 - Also **REAL**, **DOUBLE (DOUBLE PRECISION)**
 - DEC(i,j)** , **DECIMAL(i,j)** or **NUMERIC(i,j)**: decimal number
- **Character string:**
 - CHARACTER(n)** or **CHAR(n)**: fixed-length character string
 - VARCHAR(n)**: varying length character string
- **Bit string:**
 - BIT(n)**: fixed-length bit string
- More types at <https://mariadb.com/kb/en/library/data-types>

2.2. Data Types (2)

- **Date and time:**

DATE: Made up of year-month-day in the format yyyy-mm-dd

TIME: Made up of hour:minute:second in the format hh:mm:ss

TIME(i): Made up of hour:minute:second plus i additional digits specifying fractions of a second.

Format is hh:mm:ss:ii...i

TIMESTAMP: Has both DATE and TIME components

2.4. CREATE TABLE Statement (1)

- The **CREATE TABLE** command is used to specify a new relation, its attributes, their data type and constraints.
- The **key, entity integrity and referential integrity constraints** can be specified within a CREATE TABLE statement or later using the ALTER TABLE command.

```
CREATE TABLE EMPLOYEE  
(FNAME          VARCHAR(15)          NOT NULL,  
  MINIT         CHAR,  
  LNAME         VARCHAR(15)         NOT NULL,  
  SSN           CHAR(9)             NOT NULL,  
  BDATE        DATE,  
  ADDRESS       VARCHAR(30),  
  SEX           CHAR,  
  SALARY        DECIMAL(10,2),  
  SUPERSSN      CHAR(9),  
  DNO           INTEGER             NOT NULL,  
  PRIMARY KEY (SSN),  
  FOREIGN KEY (DNO) REFERENCES DEPARTMENT(DNUMBER),  
  FOREIGN KEY (SUPERSSN) REFERENCES EMPLOYEE(SSN) );
```

2.4. CREATE TABLE Statement (2)

- An example with attribute **default values** and **secondary key specification**.

```
CREATE TABLE DEPARTMENT  
(DNAME VARCHAR(10) NOT NULL,  
DNUMBER INT NOT NULL,  
MGRSSN CHAR(9) NOT NULL DEFAULT  
'888665555',  
MGRSTARTDATE DATE,  
PRIMARY KEY (DNUMBER),  
UNIQUE (DNAME),  
FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE(SSN) );
```

2.4. CREATE TABLE Statement (3)

- A constraint can be given a name following the keyword **CONSTRAINT**. Giving names to constraints is optional but is a good practice to follow.

The names of all constraints within a particular schema must be **unique**.

A constraint name can be used to **drop a constraint** (and replace it later by another constraint).

- SQL has a default policy that *any modifications violating the referential integrity constraints is rejected by the system*. However, in SQL the schema designer can specify the **action to be taken** upon *deletion of a referenced tuple* or *upon modification of a referenced primary key value*, by attaching a **referential triggered action clause** to a foreign key constraint.

The options include **SET NULL**, **CASCADE**, **SET DEFAULT**.

They can be qualified on either **ON DELETE** or **ON UPDATE**.

2.4. CREATE TABLE Statement (4)

- Example.

CREATE TABLE EMPLOYEE

(FNAME VARCHAR(15) NOT NULL,

...

DNO INTEGER NOT NULL DEFAULT 4,

CONSTRAINT EMPPK

PRIMARY KEY (SSN),

CONSTRAINT EMPSUPERFK

FOREIGN KEY (SUPERSSN) REFERENCES EMPLOYEE

ON DELETE SET NULL ON UPDATE CASCADE

CONSTRAINT EMPDEPTFK

FOREIGN KEY (DNO) REFERENCES DEPARTMENT

ON DELETE SET DEFAULT ON UPDATE CASCADE);

2.4. CREATE TABLE Statement (5)

- Example.

```
CREATE TABLE DEPARTMENT
(DNAME  VARCHAR(10) NOT NULL,
DNUMBER INT NOT NULL,
MGRSSN CHAR(9) NOT NULL DEFAULT '888665555',
MGRSTARTDATE DATE,

CONSTRAINT DEPTPK
PRIMARY KEY (DNUMBER),

CONSTRAINT DEPTSK
UNIQUE (DNAME),

CONSTRAINT DEPTMGRFK
FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE
ON DELETE SET DEFAULT ON UPDATE CASCADE);
```

- The name of a schema can be explicitly specified in the definition of a table:

```
CREATE TABLE COMPANY.DEPARTMENT ...
```

2.5. DROP DATABASE Command

- A schema can be deleted using the **DROP DATABASE** command.
- IF EXISTS can be used to prevent an error if the database doesn't exist.
- **USE THIS COMMAND VERY CAREFULLY**

DROP SCHEMA IF EXISTS COMPANY ;

2.6. DROP TABLE Command

- A table can be deleted using the **DROP TABLE** command.
- IF EXISTS can be used to prevent errors
- If a foreign key references this table, the table cannot be dropped.
In this case, it is necessary to drop the foreign key first.
- **USE THIS COMMAND WITH EXTREME CAUTION**

DROP TABLE IF EXISTS DEPENDENT;

2.7. ALTER TABLE Command (1)

- Adding a column.

```
ALTER TABLE COMPANY.EMPLOYEE ADD JOB VARCHAR(12);
```

```
ALTER TABLE COMPANY.EMPLOYEE ADD JOB VARCHAR(12) NOT  
NULL;
```

A **DEFAULT** clause can also be specified. If no such clause is specified, the new attribute will have NULLs in all the tuples of the relation after the command is executed.

- Dropping a column.

```
ALTER TABLE COMPANY.EMPLOYEE DROP ADDRESS CASCADE;
```

CASCADE option: all constraints and views that reference the column are dropped from the schema.

RESTRICT option: the command is successful only if no views or constraints reference the column.

2.7. ALTER TABLE Command (2)

- Dropping a default clause in a column.

```
ALTER TABLE COMPANY.DEPARTMENT ALTER MGRSSN DROP  
DEFAULT;
```

- Adding a default clause in a column.

```
ALTER TABLE COMPANY.DEPARTMENT ALTER MGRSSN SET  
DEFAULT '33344555';
```

2.7. ALTER TABLE Command (3)

- Dropping a constraint.

To be dropped **a constraint must have a name.**

```
ALTER TABLE COMPANY.EMPLOYEE DROP CONSTRAINT  
EMPSUPERFK CASCADE;
```

CASCADE and RESTRICT have meaning only if the constraint is a candidate key definition. In this case, if CASCADE is specified, all the foreign key definitions that reference the candidate key will be dropped too.

- IF NOT EXISTS can be used to prevent errors.

- Adding a constraint.

```
ALTER TABLE COMPANY.EMPLOYEE ADD CONSTRAINT  
EMPSUPERFK FOREIGN KEY EMPLOYEE(SUPERSSN) REFERENCES  
EMPLOYEE(SSN);
```

- The constraint to be created can be named or unnamed.

2.8. CIRCULAR FOREIGN KEYS (1)

the chicken-and-egg problem

- Defining foreign keys in tables raises some other issues:
- Consider the table definitions:

CREATE TABLE EMPLOYEE

...

FOREIGN KEY (DNO) REFERENCES DEPARTMENT(DNUMBER)

and

CREATE TABLE DEPARTMENT

...

FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE(SSN)

- The problem is that either EMPLOYEE or DEPARTMENT table has to be defined first.

In both cases the system will issue an error message.

Why?

2.8. CIRCULAR FOREIGN KEYS (2)

the chicken-and-egg problem

- One solution is to postpone the introduction of the foreign-key constraint in the first table.
- For instance, if `CREATE TABLE EMPLOYEE` is executed first, we should not have the foreign key clause `FOREIGN KEY (DNO) ...` in the statement.
- After `CREATE TABLE DEPARTMENT` has been processed, we can add the desired constraint to employee using the `ALTER TABLE` command:

```
ALTER TABLE EMPLOYEE  
ADD CONSTRAINT EMPDEPTFK  
FOREIGN KEY (DNO) REFERENCES  
DEPARTMENT(DNUMBER)
```

2.8. CIRCULAR FOREIGN KEYS (3)

the chicken-and-egg problem

- If we want to populate our COMPANY database, we are in front of another problem.
- If we try to insert a tuple to relation EMPLOYEE with any non-NULL value for attribute DNO, *our insertion will be rejected as it violates the foreign key constraint*
DNO references DEPARTMENT(DNUMBER).

If we try to insert a tuple to relation DEPARTMENT with any non-NULL value for attribute MGRSSN, *our insertion will be rejected as it violates the foreign key constraint*
MGRSSN references EMPLOYEE(SSN).

- We are again facing the chicken-and-egg problem!

2.8. CIRCULAR FOREIGN KEYS (4)

the chicken-and-egg problem

- One solution is to initially replace DNO component in all tuples to be inserted in the EMPLOYEE relation with NULL. Then, when DEPARTMENT is populated with appropriate tuples, we can scan relation EMPLOYEE and replace NULLs with valid department numbers.
- This solution is *awkward* and *error-prone*.
- A better solution is to use *transactions* and *deferred checking of integrity constraints*. We will discuss this issue later (if we have time).